




# Dynamic & Adaptive Inference Tuning (Reinforcement Learning & Other Methods)

*By Chris Fregly*



<https://www.amazon.com/Systems-Performance-Engineering-Optimizing-Inference/dp/B0F47689K8/>


# Why focus on inference tuning?

Faster time-to-response

Better end-user experience

Serve more users with same hardware

GPU cost savings (\$\$\$ => \$)



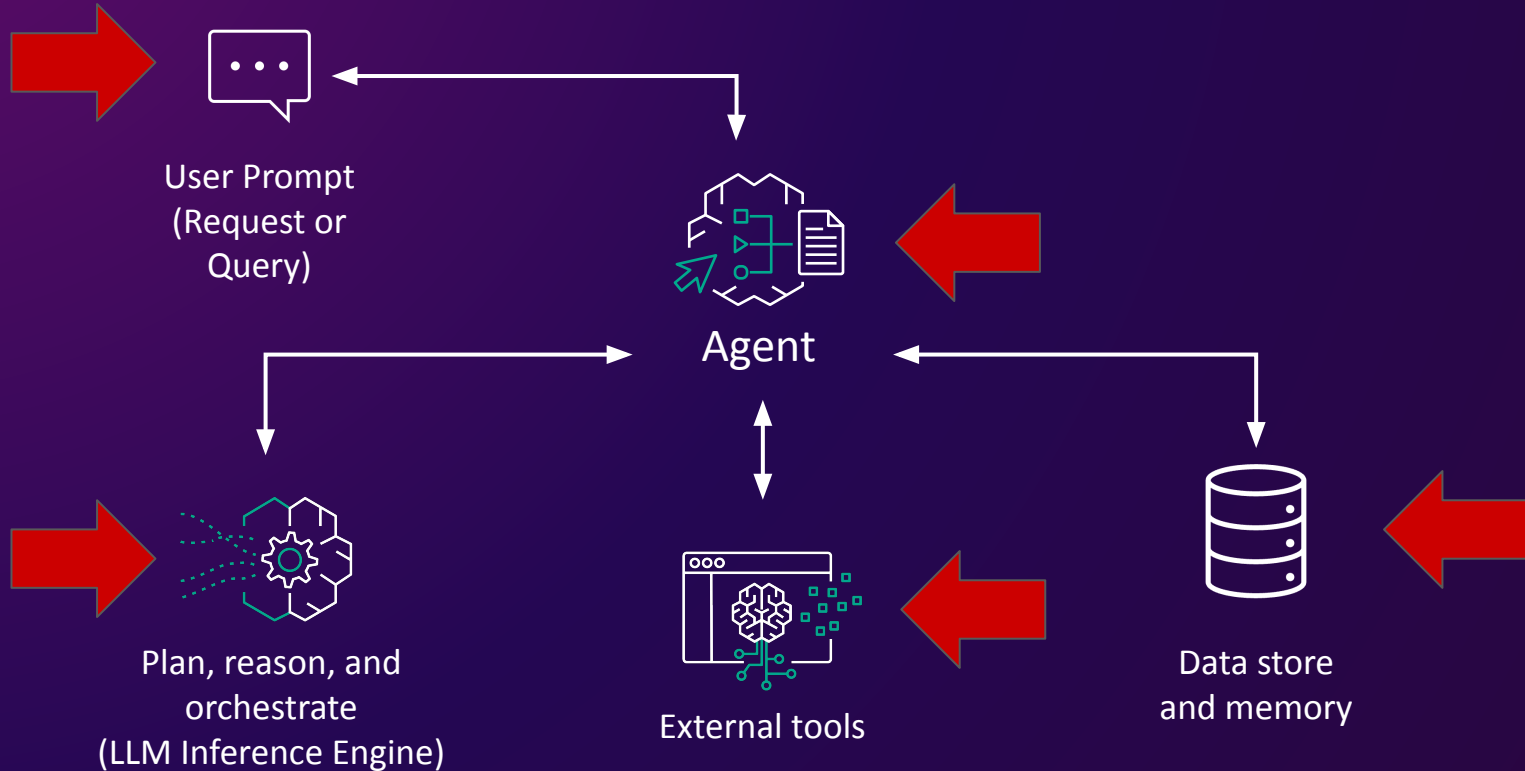
# "Mechanical sympathy"

**Original reference:** Race car driver, Jackie Stewart, who was famously aware of the inner workings his race car.


-- Martin Thompson (<https://mechanical-sympathy.blogspot.com>)

**Today's computing reference:** Co-designing software and algorithms hand-in-hand with hardware capabilities to maximize performance.

# Where are the bottlenecks?



# Where are the bottlenecks that *I will cover today?*



# Agenda

- 1 Dynamic and Adaptive AI System Optimizations
- 2 RL Agents for Production Systems Tuning and Operations
- 3 AI-Assisted Kernel and System Tuning

# Dynamic and Adaptive AI System Optimizations

## Adaptive Parallelism Switching

=>

Inference Traffic Pattern	Recommended Parallelism	Rationale
Many short requests (<256 tokens, high QPS)	Data Parallel / Replica Scaling	Minimizes inter-GPU communications, each GPU runs replicas handling independent requests (assuming the model fits into a single GPU's memory)
Few long requests (≥8k tokens, low concurrency)	Pipeline Parallelism (with micro-batches)	Reduces per-request latency by splitting layers across GPUs
Mixed load (short + some long)	Hybrid Dynamic (auto-switching)	Runs small chats on single GPUs, pipelines long ones to meet latency SLAs
Extremely large model (>1 GPU memory)	Tensor + Pipeline Hybrid	Required to fit model; balances compute and memory across both dimensions
MoE model inference (sparse expert selection)	Expert Parallelism	Distributes individual experts across GPUs; each request only invokes a subset of experts, reducing per-device memory and compute load

## Dynamic Precision Changes

## Adaptive Batching and Scheduling (Chunked Prefill configuration, etc.)

## Adjustable Speculative Decoding and KV-Cache Prefetching Configuration (Draft Model)

## Occupancy-Aware CUDA Kernel Launch Parameters

## Hot-Swappable CUDA Kernel Implementations =>

```
Python
import new_flash_attn_lib

# Monkey-patch the model's attention forward to use the new library
old_attn_forward = model.transformer.self_attn.forward

def new_attn_forward(self, *args, **kwargs):
    return new_flash_attn_lib.forward(*args, **kwargs)

model.transformer.self_attn.forward =
    new_attn_forward.__get__(
        model.transformer.self_attn,
        type(model.transformer.self_attn))
```

# RL Agents for Production Systems Tuning and Operations

## RL-based, Adaptive System Tuning

=>


**Action 1:** choose parallelism mode: single, TP, PP, and hybrid  
**Action 2:** choose precision: full FP8 versus mixing FP8 and FP4  
**Action 3:** adjust batch size or batch-waiting time  
**Action 4:** enable or disable cache compression  
**Action 5:** enable or disable speculative decoding  
**Action 6:** select a smaller draft model for speculative decoding  
**Action 7:** select a larger draft model for speculative decoding  
**Action 8:** enable or disable speculative kv prefetching

## Congestion-Aware Disaggregated Prefill-Decode Request Routing

## RL Agent for Real-Time Monitoring, Ops, and Issue Resolution

## RL-based Mixture-of-Experts (MoE) Expert Routing


## GPU-aware Prefill-Decode Kernel Placement =>






# AI-Assisted Kernel and System Tuning


**Alpha Tensor Project (Google DeepMind).** Rediscovered Strassen's sub-quadratic GEMM algorithm. Optimized CUDA kernel (better than NVIDIA).



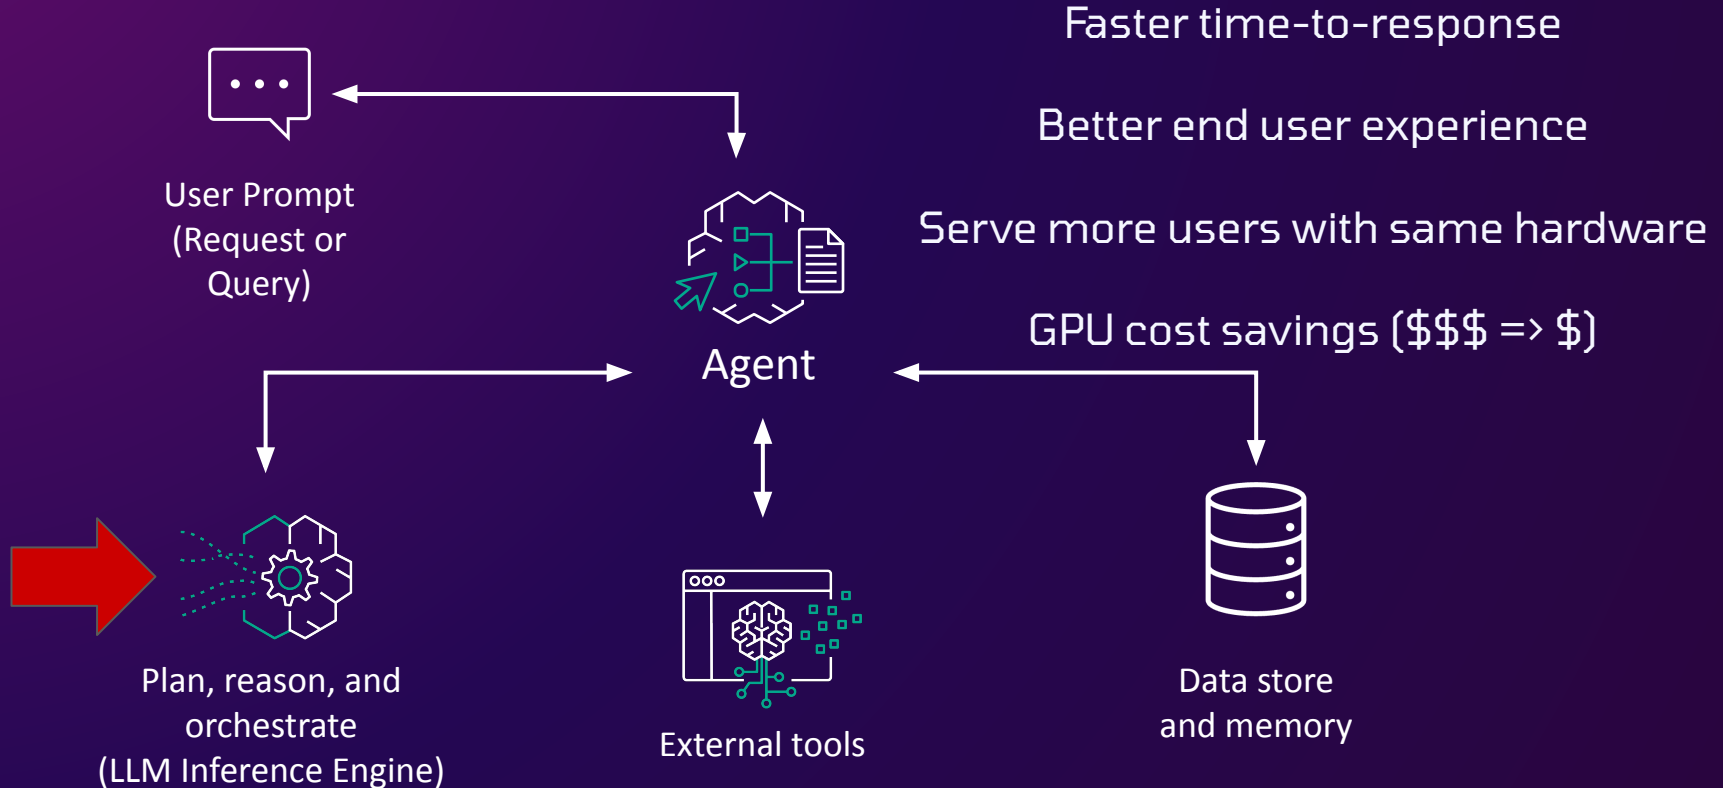
**Optimize Attention CUDA kernel (NVIDIA).** Used DeepSeek-R1 reasoning model to create optimized Attention CUDA kernel implementation.



**"AI CUDA Engineer" (Sakana.ai).** Evolutionary strategies to iteratively refine CUDA kernel code. Optimized CUDA kernel (5x NVIDIA performance). Also, was able to process 230 of 250 primitive PyTorch ops.


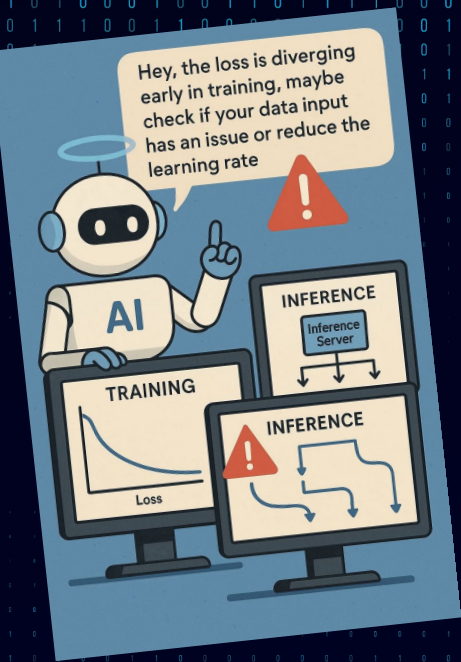


# Embrace Adaptive, AI-Assisted System Optimizations



# Thank You!

## Dynamic & Adaptive Inference Tuning (Reinforcement Learning & Other Methods)



*By Chris Fregly*

<https://www.amazon.com/Systems-Performance-Engineering-Optimizing-Inference/dp/B0F47689K8/>